

Code Implementation description

Setup & Installation (Colab Environment)

1. Library Imports:

- Imports necessary Python libraries such as os, shutil, PIL (for image processing), torch (PyTorch), and torchvision's optical_flow models.
- Includes optional Colab-specific tools like files for uploading/downloading.

2. Change Working Directory:

- Ensures the script operates from the /content directory, the default workspace in Google Colab.

3. Cloning GFPGAN Repository:

- Removes any existing GFPGAN directory to avoid conflicts.
- Clones the official [GFPGAN repository](#) from GitHub, which provides tools for face restoration and enhancement.

4. Installing Required Dependencies:

- Uses pip to install:
 - numpy==1.24.4 (compatible version).
 - Required deep learning libraries: basicsr, facexlib, realesrgan.
 - Dependencies listed in GFPGAN's requirements.txt.
- Runs python setup.py develop to install GFPGAN in development mode (enables live code updates without reinstalling).

5. Compatibility Fix for torchvision:

- Fixes a deprecated import statement in the installed basicsr package that is incompatible with recent versions of torchvision (e.g., Python 3.11).
- Uses sed to update the import of rgb_to_grayscale to the correct path.

6. Download Pre-trained GFPGAN Model:

- Creates a directory experiments/pretrained_models to store model weights.
- Downloads the GFPGANv1.3.pth pre-trained model for inference use.

GIF Upload and Frame Extraction

1. Define Input Directory:

- `gif_upload = 'inputs/upload'`: Specifies the directory where extracted frames will be saved.
- `shutil.rmtree(...)`: Removes any existing content in the directory to start clean.
- `os.makedirs(...)`: Recreates the directory to store new frames.

2. Upload GIF File (Google Colab):

- `files.upload()`: Opens a file picker for the user to upload a GIF.
- `gif_file = list(uploaded.keys())[0]`: Retrieves the uploaded filename from the returned dictionary.

3. Extract Frames from the GIF:

- `Image.open(gif_file)`: Opens the uploaded GIF using PIL.
- `ImageSequence.Iterator(img)`: Iterates through each frame of the animated GIF.
- `frame.convert("RGB").save(...)`: Converts each frame to RGB format and saves it as a PNG file in the `gif_upload` directory.
- `frame_paths.append(...)`: Tracks each saved frame's file path.

4. Progress Output:

- Prints the number of frames extracted from the uploaded GIF, confirming successful preprocessing.

✦ Face Restoration with GFPGAN

1. Clear Previous Results:

- `!rm -rf results`: Deletes any previous output directory to avoid mixing old and new results.

2. Run Inference with GFPGAN:

- `inference_gfpgan.py` is the main inference script from the GFPGAN repository.
- Parameters:
 - `-i {gif_upload}`: Input folder containing extracted frames from the GIF.
 - `-o results`: Output directory where enhanced images will be saved.
 - `-v 1.3`: Uses the GFPGAN version 1.3 model (`GFPGANv1.3.pth`).
 - `-s 2`: Sets the upscale factor to 2×.
 - `--bg_upsampler realesrgan`: Uses Real-ESRGAN for background upscaling to improve image quality beyond faces.

3. Collect Restored Frame Paths:

- `restored_dir = 'results/restored_imgs'`: Points to the directory where enhanced frames are saved.
- `restored_paths = sorted(...)`: Creates an ordered list of all restored frame file paths with `.png` extension.

Motion Estimation Using RAFT Optical Flow

1. RAFT Model Setup:

- Uses `raft_large` model from `torchvision.models.optical_flow`, with pretrained weights (`Raft_Large_Weights.DEFAULT`).
- Moves the model to GPU if available and sets it to evaluation mode (`.eval()`).

2. Image Preprocessing:

- `load_frame(...)`: Opens each restored image, resizes it to (384, 384) for consistent input shape, and converts it into a normalized PyTorch tensor with batch dimension [1, 3, H, W].

3. Preparing Input Tensors:

- Loads and stores tensors for each restored frame from `restored_paths` using the `load_frame()` function.

4. Debugging & Device Info:

- Prints whether GPU is used and displays `PYTORCH_CUDA_ALLOC_CONF` for potential debugging of memory issues.

5. Optical Flow Computation:

- Iterates over each consecutive frame pair.
- Applies the RAFT model's internal transform pipeline for proper preprocessing (e.g., normalization).
- Inference is wrapped with `torch.no_grad()` to avoid gradient tracking and save memory.
- Only the final refined flow result from RAFT is used.
- Estimated flow tensors are moved to CPU to free up GPU memory.

6. Output:

- Prints a success message indicating how many frame pairs had their motion vectors estimated.

Image Warping and Temporal Loss Calculation

This section aims to:

- **Warp** each frame toward the next frame using estimated optical flow.
 - **Quantify temporal consistency** using Mean Squared Error (MSE) between warped and actual frames.
-

Implementation Details

1. Setup:

- Initializes warped list to store warped frames and losses to store temporal losses.
- Uses `torch.nn.MSELoss()` as the criterion for frame difference measurement.

2. Per-Frame Warping Loop:

For each frame pair:

- Moves both input (`img`) and target (`target`) frames to the computation device (GPU or CPU).
- Moves optical flow (`flow`) to the same device (fixes mismatch errors).
- Generates a 2D mesh grid (`grid_x`, `grid_y`) representing pixel coordinates.
- Combines these into a grid and adds the flow to warp it: `vgrid = grid + flow`.

3. Normalization for Grid Sampling:

- RAFT outputs flow in pixel units. `F.grid_sample` requires normalized coordinates in the range `[-1, 1]`.
- The code scales the `vgrid` accordingly and permutes its shape to match the input format `[B, H, W, 2]`.

4. Image Warping:

- `F.grid_sample(...)` warps the current frame (`img`) toward the next frame using the transformed grid.
- This simulates how the current frame would look if moved using the predicted motion (optical flow).

5. Temporal Loss Calculation:

- Calculates MSE between the warped image and the actual next frame (target).
- Stores the warped image and its corresponding temporal loss.

6. Output:

- Prints per-frame temporal loss values, which indicate how well the motion warping aligns with the next frame. Lower values imply better temporal coherence.

GIF Reconstruction from Enhanced Frames

Purpose:

After performing face restoration and optional motion analysis (optical flow), this step reassembles the enhanced individual frames into a single animated GIF.

Implementation Details

1. Load Enhanced Frames:

- `restored_images = [Image.open(p).convert("RGB") for p in restored_paths]:`
 - Opens each restored PNG file.
 - Ensures all frames are in RGB format for consistency.

2. Define Output Path:

- `output = '/content/outputs/enhanced_output.gif':` Specifies where the final GIF will be saved.
- `os.makedirs(...):` Ensures the output directory exists.

3. Save as Animated GIF:

- `restored_images[0].save(...):` Saves the first frame and appends the rest using:
 - `save_all=True:` Enables multiple frame saving.
 - `append_images=...:` Adds the remaining frames.

- `duration=100`: Sets frame duration (in milliseconds) — 10 frames per second.
- `loop=0`: Makes the animation loop infinitely.

4. Output Confirmation:

- Prints the path where the enhanced GIF was saved.

Objective Quality Assessment (PSNR & SSIM)

Purpose:

To compare the **restored (enhanced)** frames against the **original low-quality** frames and **quantify visual improvements** using widely accepted image quality metrics:

- **PSNR (Peak Signal-to-Noise Ratio)**: Measures the ratio between the maximum possible pixel value and the power of the noise (difference between the images). Higher is better.
- **SSIM (Structural Similarity Index)**: Evaluates perceived quality by considering luminance, contrast, and structural changes. Ranges from -1 to 1, where **1 indicates perfect similarity**.

Implementation Details

1. Frame Loading & Alignment:

- Loads each original and restored frame from:
 - `inputs/upload`: Original extracted frames.
 - `results/restored_imgs`: Enhanced frames after GFPGAN processing.
- Resizes both to **384×384** pixels for consistent comparison.

2. Metric Calculation:

- **PSNR** is calculated using `skimage.metrics.peak_signal_noise_ratio`.
- **SSIM** is calculated using `skimage.metrics.structural_similarity` with `channel_axis=2` for RGB images.
- Metrics are computed for each frame pair and stored in lists.

3. Averaged Output:

- The script prints the **mean PSNR** and **mean SSIM** across all frames:

python

CopyEdit

✓ PSNR (avg): 32.54 dB

✓ SSIM (avg): 0.9124

✦ Perceptual Quality & Temporal Consistency Evaluation

To assess the quality of the restored video frames, we used two important evaluation methods:

✓ 1. Temporal Loss (Smoothness Across Frames)

- **Purpose:** Measures how smoothly frames transition over time in the restored video.
- **Why it matters:** Flickering or inconsistent frames can reduce visual quality. A lower temporal loss means the restoration is temporally stable and pleasant to watch.
- **Output:** An average score showing how consistent the restored frames are. Lower is better.

✓ 2. LPIPS – Perceptual Image Quality

- **Purpose:** LPIPS (Learned Perceptual Image Patch Similarity) evaluates how *visually similar* the restored images are compared to the original ones.
- **Why it matters:** Traditional metrics like PSNR or SSIM don't always reflect human visual perception. LPIPS uses deep neural networks to model how people actually see differences.
- **Process:** Each restored frame is compared to its original version. The LPIPS model outputs a perceptual difference score for each pair.
- **Output:** An average LPIPS score across all frames. Lower values indicate that the restored images are closer to the original in terms of perceptual quality.